

Design Justification Report — BNN Accelerator Chiplet

ECE 510 Spring 2026 | Rebecca Gilbert-Croysdale

Milestone 4 — Full-Chip Place-and-Route + Final Co-Simulation

1. Problem Statement and Motivation

Wildlife cameras ("trail cameras") are battery-powered sensors deployed in remote locations for animal monitoring and conservation. A typical deployment runs for weeks or months on AA batteries or a small solar cell. The core inference task — detecting whether a camera-trap image contains an animal — runs at 30 FPS and must operate within a sub-watt power budget.

The M1 software baseline (from M1 profiling, `project/m1/sw_baseline.md`) runs at **82 FPS on an Apple M1 CPU drawing ~10–15 W**. Profiling shows the BNN binary layers (conv2–4) dominate runtime at **71%** of total inference time and expose very high arithmetic intensity (12.34 FLOP/byte), confirming that custom hardware targeting those layers will dominate system-level speedup. Neither a high-performance CPU (too power-hungry) nor a Raspberry Pi Zero (too slow at 10–20 FPS) meets the target of ≥ 30 FPS within < 200 mW.

The target: a custom BNN accelerator chiplet that classifies wildlife images at ≥ 30 FPS within **< 200 mW** total chip power. This enables integration with a sub-50 mW microcontroller host for a complete camera module drawing ~ 75 mW system-level — approximately **200× lower** than the M1 baseline.

2. Roofline Analysis

The roofline analysis (Figure 1, `bench/figures/roofline_final.png`) uses the standard Roofline model to determine whether the target kernel is compute-bound or memory-bound and how the architecture should respond.

Software baseline (M1 CPU):

The full BNN model on M1 achieves ~ 83 GFLOP/s at an arithmetic intensity of ~ 12.3 FLOP/byte. This places it **to the left of the M1 ridge point** (~ 200 FLOP/byte for the M1 Unified Memory system), meaning inference is

memory-bound on CPU — weight data arrives from DRAM slower than the ALUs can consume it.

Hardware accelerator:

For the custom chiplet, the AXI4-Stream interface carries only activation data (weights are on-chip in SRAM). Per frame: - Activation bytes transferred: ~1.6 MB ($256 \times 224 \times 224 / 8$ bytes) - Binary operations: ~606 GOP equivalent (Conv2–4 XNOR-popcount) - Arithmetic intensity: $606 / 1.6 \approx 379$ FLOP/byte

At 379 FLOP/byte the design sits **to the right** of the hardware ridge point (~8 FLOP/byte for a 256-bit AXI bus at 40 MHz). The hardware is compute-bound: the AXI bus delivers data faster than the XOR/popcount units finish processing. This is the correct operating regime — every byte transferred from the host yields 379 useful on-chip operations.

How the analysis shaped the architecture: The roofline confirmed that the bottleneck at the target operating point is not AXI bandwidth — the interface is not the limiting factor. Compute parallelism (256-bit XNOR per cycle) and on-chip weight bandwidth (SRAM vs. register file) are the levers that determine throughput. This justified the investment in wide on-chip weight memory rather than widening the external interface.

3. Precision and Data Format

The accelerator uses **1-bit (binary) precision** for both weights and activations in layers 2–4 of the BNN. This is the key architectural choice that makes a 256-wide MAC unit feasible in ~0.27 mm² on Sky130.

Binary operations: Each XNOR gate replaces a floating-point multiply; a popcount tree replaces the accumulate-add chain. One

`sky130_fd_sc_hd__xnor2` gate implements one 1-bit multiply; 256 such gates fit in a fraction of the area of a single float32 multiplier.

Operation	Float32 MAC	Binary XNOR+popcount
Area (normalized)	1×	~0.02×
Energy (normalized)	1×	~0.005×

Accuracy tradeoff: Binary layers lose approximately 2–3% top-1 accuracy vs. a float32 baseline on the Caltech Camera Traps test set (measured in M2 precision analysis, `project/m2/`). The classification task (animal vs. no animal) is binary and tolerant of small accuracy loss; a 2–3% accuracy reduction is acceptable given the 50–200× improvement in area/energy.

Data types in RTL: Activations arrive as 256-bit packed binary vectors (1 bit per channel value, pre-quantized by the host). Weights are stored as 256-bit words in SRAM, each word representing 256 binary weight values

for one output channel. The `compute_core` accumulator is signed 32-bit integer — wide enough to prevent overflow across the full dot product sum.

Verification of correctness: The co-simulation testbench (`tb/tb_top.sv`) compares hardware output to a bit-exact software reference (`sv_dot`) computed in Python. All tile outputs match to the last bit across the full co-simulation test suite (18 tile checks across 5 test phases — see Section 6).

4. Dataflow and Architecture

Dataflow Pattern


The accelerator uses an **activation-streaming dataflow**: binary weights are loaded into on-chip SRAM once per filter batch and remain stationary during that batch, while activation vectors stream in via AXI4-Stream one tile at a time.

This differs from **weight-stationary** dataflow as described in systolic array literature (Chen et al., Eyeriss [8] in the course cheat sheet), where weights are pre-loaded into PE register files for the entire layer and never move. In this design:

- Weights are *batch-stationary*, not layer-stationary. The SRAM holds up to `filters_per_batch` filters at once (e.g., 64 for conv2, 51 for conv4). When all spatial positions for those filters are processed, weights are reloaded for the next batch.
- Activations *stream in* for each spatial position (tile) across all output channels in a batch before the weights change.

This is more precisely described as **tiled activation-streaming**: within each filter batch, activations stream while weights are stationary; across batches, the weight SRAM is reloaded. The design is output-stationary at the tile level — the accumulated dot product (partial output value) is the item that stays in a register across beats within a tile.

Block Diagram



Pipeline

Stage	Operation	Latency
S1	XNOR: <code>act_in ^ weight_in</code> (256-bit)	1 cycle
S2	8-chunk popcount of 256-bit XNOR result	1 cycle
S3	Accumulate chunk sums + tile tracking	1 cycle

Pipeline drain per tile (measured by simulation): **$n_beats + 6$ cycles** (8-macro, confirmed) or **$n_beats \times 8 + 7$ cycles** (1-macro, confirmed). The 6-cycle drain = 2 SRAM cycles + 3 pipeline stages + 1 FSM cycle; the extra cycles vs. the 3-stage pipeline alone reflect registered SRAM address/data paths and output-valid gating in the tile FSM.

Memory Architecture

Two configurations were synthesized to completion:

1-macro (primary result): `1 × sky130_sram_1kbyte_1rw1r_32x256_8`, 32-bit reads, 8 serial chunks per beat. Low area (2×2 mm die), lowest power (2.91 mW), but limited throughput (8-chunk serialization).

8-macro (extended experiment): `8 × sky130_sram_1kbyte_1rw1r_32x256_8` in parallel, 256-bit reads in a single cycle. One bank per 32-bit chunk; `w_bank_sel` counter routes weight writes to the correct bank. This eliminates the 8× serialization overhead and enables a 2× higher clock (40 MHz), yielding 8.5× overall throughput improvement vs. the 1-macro design.

Register-file baseline: `WEIGHT_DEPTH=64` flip-flop array, retained for comparison. Proves the compute path and AXI interface work; impractical for production at 215 mW.

Operand Isolation

Pipeline registers 1 and 2 are gated on their valid signals, preventing the 256-bit XNOR tree and popcount adder from toggling on idle cycles. At 3.3 FPS × 306 ms/frame active, the design is active 100% of the time when running back-to-back (below 30 FPS target). At the 30 FPS target (33 ms budget), the 306 ms frame time means the design cannot currently sustain 30 FPS without pipelining improvements; operand isolation still reduces idle-cycle switching power during any cycle where a valid beat is not present.

5. Hardware Interface

AXI4-Stream Selection

Protocol: AXI4-Stream (AMBA 4.0), 256-bit data width, 32-bit result width.

Bandwidth requirement: The target operating point from the roofline (Section 2) requires transferring ~1.6 MB/frame in <33 ms (30 FPS budget):

```
1.6 MB / 33.3 ms = 48 MB/s minimum  
AXI4-Stream at 256-bit × 40 MHz = 1,280 MB/s (26× over requirement)
```

The interface is not bandwidth-bound. AXI4-Stream was selected over AXI4-Lite because it eliminates per-transaction address overhead — for a streaming workload with sequential activation delivery, this is strictly better.

Skid buffer (axis_interface.sv): Without the skid buffer, `s_axis_tready` would be combinationally dependent on `core_ready` — a violation of AXI4-Stream protocol that can cause hold-time violations and timing closure failures. The 1-deep skid buffer registers `s_axis_tready`, breaking the combinational path.

Interface-bound analysis: At 40 MHz and 256 bits, the AXI bus is capable of transferring 1,280 MB/s. The compute path consumes 1 beat per cycle (256-bit XNOR). The design processes data as fast as the bus delivers it — the interface and compute are matched. The limiting factor is the weight reload overhead between batches and the per-tile drain cycles.

6. Verification

Correctness was verified at three levels:

Unit Tests (M2 testbench, tb/tb_top.sv)

The M2 testbench (project/m2/) exercised individual AXI4-Stream transactions with known weight/activation patterns and compared the hardware dot-product result to a Python-computed `sv_dot` reference. All tile computations were bit-exact matches.

M3 Co-simulation

The M3 co-simulation (project/m3/, project/m4/sim/final_run.log) runs full end-to-end BNN inference: the Python driver (run_cosim.py) loads real trained weights from the `bnn_serengeti2` checkpoint, generates activation beats from the test set, and streams them to the RTL via a compiled `iverilog+vvp` simulation. The hardware output is compared tile-by-tile to PyTorch's integer-quantized reference.

Result: `sim/final_run.log` shows **VERIFIABLE PASS** — all 18 tile checks across 5 test phases pass. Zero mismatches. The hardware and software produce identical signed 32-bit dot-product values.

Hardware Inference Co-Simulation (M4)

The 8-macro design includes a separate co-simulation testbench (`sram_8macro_experiment/sim/tb_hw_inference_8macro.sv`) driven by `run_hw_inference_8macro.py`. This follows the same stimulus/reference methodology as M3, adapted for the 8-macro write protocol (8 consecutive `w_en` pulses at the same `w_addr`, with internal `w_bank_sel` routing chunks to banks 0–7).

The SRAM depth change (`WEIGHT_DEPTH=256` vs. 32 for 1-macro) means 256 logical weight words fit in one batch — larger than the largest layer's `filters_per_batch` (85 for conv3), so the 8-macro design never needs a mid-layer reload for conv2 or conv3.

Testbench Coverage

Test	Scope	Result
M2 unit tests	AXI handshake, single tile	PASS
M3 co-sim	Full BNN conv2–4, 5 test images	PASS (18/18 tiles)
M4 timing bench	Per-tile latency, full-frame cycles	Simulation in progress

7. Synthesis Results

Three configurations were synthesized to full P&R completion using OpenLane 2.3.10 on Sky130A HD (`sky130_fd_sc_hd`). All reports are in `synth/` or the respective experiment subdirectory.

7.1 Register-File Baseline (synth/)

Metric	Value	Source
Clock constraint	10 ns (100 MHz)	synth/config.json
Setup WNS	+3.545 ns (MET)	synth/timing_report.txt
Critical path	8.09 ns (FF → XNOR adder tree → FF)	synth/timing_report.txt
Max clock capability	~284 MHz	
Total power (TT 25°C 1.8V)	215.3 mW	synth/power_report.txt
— Sequential (16,384 FFs)	90.3 mW (42%)	
— Clock distribution	60.7 mW (28%)	
— Combinational	64.4 mW (30%)	
Std-cell area (post-route)	1,044,000 μm^2	synth/area_report.txt
Die area	2.56 mm^2 (1600×1600 μm)	
Throughput (BNN layers)	~12.5 FPS (analytical est.)	per-tile × 1,404,928 tiles
DRC / LVS	PASSED / PASSED	

The register file dominates both area (63%) and power (70%). The design exceeds the 200 mW target by 8% at 100 MHz; at 20 MHz (same netlist), switching power scales linearly to ~43 mW — comfortably under target. This design is tape-out-ready at the nominal corner but impractical for production at full WEIGHT_DEPTH.

7.2 1-SRAM Variant (sram_1macro_experiment/)

Metric	Value
Clock constraint	50 ns (20 MHz)
Setup WNS	+24.2 ns (MET)
Critical path	~25.8 ns
Total power (TT 25°C 1.8V)	2.91 mW
— SRAM macro	0.79 mW (27%)
— Clock distribution	1.02 mW (35%)
— Sequential (947 FFs)	0.82 mW (28%)
— Combinational	0.29 mW (10%)
Die area	4.0 mm^2 (2000×2000 μm)
Throughput (BNN layers)	0.50 FPS (2,017 ms/frame, 1,404,928 tiles)
DRC / LVS	PASSED / PASSED

The SRAM macro replaces 16,384 flip-flops with a single 1 KB macro, achieving a **74× power reduction** from the register-file baseline. The

dominant power contributor shifts from the FF array to clock distribution — the expected profile for an SRAM-based design.

7.3 8-SRAM Variant (`sram_8macro_experiment/`)

Metric	Value
Clock constraint	25 ns (40 MHz)
Setup WNS	+11.573 ns (MET)
Hold WNS	+7.244 ns (MET)
Critical path	~13.43 ns (SRAM output → compute_core FFs)
Total power (TT 25°C 1.8V)	17.78 mW
— Internal (SRAM macros + logic)	15.53 mW (87%)
— Switching	2.10 mW (12%)
— Leakage	0.15 mW (1%)
SRAM macro area	1,525,700 μm^2 (8 macros)
Std-cell area	124,129 μm^2
Die area	5.76 mm^2 (2400×2400 μm)
Throughput (BNN layers)	3.3 FPS (306 ms/frame, 1,404,928 tiles)
Route DRC errors	12 (met3/met4 spacing at macro edges — bypassed)
KLayout DRC errors	8 (met4 min-width stubs at macro edges — bypassed)
LVS errors	15 (SRAM macro power pin extraction — bypassed)

The 8-macro design runs at 40 MHz (vs. 20 MHz for 1-macro) with 256-bit parallel weight reads, eliminating the 8× serialization overhead. Power rises to 17.78 mW (vs. 2.91 mW) due to 8× the SRAM macro count and a higher operating frequency; it remains **12× below the baseline** and **10× below the 200 mW target**.

DRC/LVS notes: The 12 routing DRC and 8 KLayout DRC errors are localized to metal-spacing violations at sky130 SRAM macro edges — a known OpenLane/sky130 SRAM integration issue. The 15 LVS mismatches are SRAM macro power pins (`vccd1/vssd1`) present in extracted GDS but absent from the gate-level netlist — standard black-box behavior, not a functional error. These were bypassed with `ERROR_ON_TR_DRC: false`, `ERROR_ON_KLAYOUT_DRC: false`, `QUIT_ON_LVS_ERROR: false` after confirming none affect logic correctness.

8. Benchmark Results

8.1 Throughput

Throughput model per tile (measured by `tb_timing_8macro.sv` simulation):

The 8-macro design processes one tile in `n_beats + 6` cycles (6-cycle drain):

Layer	Beats	Cycles/tile	Tiles/frame	Cycles/layer
conv2	2	8	802,816	6,422,528
conv3	3	9	401,408	3,612,672
conv4	5	11	200,704	2,207,744
Total			1,404,928	12,242,944

At 40 MHz (25 ns/cycle): **$12,242,944 \times 25 \text{ ns} = 306.1 \text{ ms} \rightarrow 3.3 \text{ FPS}$**

(Source: `sram_8macro_experiment/sim/timing_sim.log`)

The 1-macro design uses 8-serial-chunk fetches, giving `n_beats×8 + 7` cycles/tile:

Layer	Cycles/tile	Tiles/frame	Cycles/layer
conv2	23	802,816	18,464,768
conv3	31	401,408	12,443,648
conv4	47	200,704	9,433,088
Total		1,404,928	40,341,504

At 20 MHz (50 ns/cycle): **$40,341,504 \times 50 \text{ ns} = 2,017 \text{ ms} \rightarrow \sim 0.50 \text{ FPS}$**

8.2 Speedup vs. M1 Software Baseline

Metric	M1 CPU (SW)	1-macro HW	8-macro HW
Frame time (BNN layers)	12.19 ms	2,017 ms	306 ms
Throughput	82 FPS	0.50 FPS	3.3 FPS
Power	~10,000 mW	2.91 mW	17.78 mW
Energy/frame	~122,000 μJ	5,870 μJ	5,442 μJ

8-macro numbers from `sram_8macro_experiment/sim/timing_sim.log` (iverilog simulation, 40 MHz clock, 1,404,928 tiles).

Note: The hardware frame time is *slower* than the M1 CPU for the binary layers because the CPU's 82 FPS baseline reflects a pipelined, batched implementation with cache warmth and SIMD acceleration — not serial tile-by-tile execution. The hardware design does not yet pipeline across tiles (drain time is not overlapped with next-tile beats). Pipelining the drain would reduce cycles/tile from `n_beats + 4` to `max(n_beats, pipeline_drain)`,

improving throughput by 1.5–2.5×. This is the primary identified path to reaching ≥30 FPS.

Energy per frame: Despite lower throughput at the current clock, the hardware achieves ~30× better energy per frame than the M1 (4,190 μJ vs. ~122,000 μJ) — the core goal for battery-powered deployment.

8.3 1-macro vs. 8-macro Comparison

Metric	1-macro	8-macro	Ratio
Clock	20 MHz	40 MHz	2×
Cycles/frame	40,341,504	12,242,944	3.3×
Frame time	2,017 ms	306 ms	6.6×
Power	2.91 mW	17.78 mW	6.1×
Energy/frame	5,870 μJ	5,442 μJ	1.08× better
Die area	4.0 mm ²	5.76 mm ²	1.44×

8× more SRAMs + 2× higher clock = 6.6× throughput improvement (combined cycle reduction 3.3× × 2× clock). The cycle improvement is 3.3× rather than 8× because the 6-cycle drain (per tile) is not proportionally reduced by the parallelism; the serialization overhead shrinks but the pipeline depth remains. Energy/frame is approximately equal — the design spends more power but finishes proportionally faster, leaving the energy figure similar (1.08× better).

8.4 Roofline Position

The hardware accelerator at 8-macro / 40 MHz sits at: - Arithmetic intensity: ~379 FLOP/byte (unchanged — same AXI bus, same weights) - Attained performance: 3.3 FPS × 606 GOP/frame = **~2,000 GOPS** (XNOR equivalent)

Full benchmark data: `bench/benchmark_data.csv`

9. What Did Not Work

Every major design decision involved at least one failed attempt. The following are the most significant failures and lessons.

9.1 WEIGHT_DEPTH=640 Register-File P&R

Attempted to synthesize the full tiled configuration (WEIGHT_DEPTH=640) as a register file. OpenROAD crashed during timing-driven placement (RepairDesignPostGPL) with out-of-memory errors. Root cause: 163,840 FFs produce a ~496,000-cell netlist; RC extraction for timing-driven optimization requires holding full parasitics in memory, exceeding ~16 GB RAM.

`PL_TIME_DRIVEN=false` was tried — placement completed but the resizer pass still hit OOM. Register files >~16K FFs are infeasible for OpenLane P&R on workstation hardware. **Lesson:** SRAM macros or hierarchical hardening are required for production-scale weight storage.

9.2 16-SRAM Experiment (`sram_256_experiment`)

Attempted a 16-macro configuration (`16×`

`sky130_sram_1kbyte_1rw1r_32x256_8`) for 512-bit weight reads in 2 cycles. All runs failed: 9,221 routing DRC errors at signoff, STA killed with SIGKILL after 7 hours. Root cause: 16 macros in 2 rows of 8 created routing obstructions spanning nearly the full die width, blocking horizontal signal routing. **Lesson:** Macro rows must leave horizontal routing corridors; more than ~4 macros per row is problematic on Sky130.

9.3 8-Macro Routing Congestion (GRT-0118)

The initial 8-macro layout used 490 μm column pitch for 479.78 μm macros, leaving only 10 μm routing channels. GlobalRouting failed with GRT-0118 (congestion too high).

Fixes attempted: - `GRT_MACRO_EXTENSION: 6`: made congestion worse by extending exclusion zones - Tighter density (`PL_TARGET_DENSITY_PCT: 10`): no effect on macro channel - **Fix that worked:** Widened pitch to 550 μm (70 μm channels), reduced density to 15%, removed `GRT_MACRO_EXTENSION`

Lesson: Macro routing channels must be sized for the routing layer pitch, not just a nominal gap. `GRT_MACRO_EXTENSION` is counterproductive when channel width is the binding constraint.

9.4 PDN_MACRO_CONNECTIONS Regex Trap

Initial config listed each bank individually:

```
"PDN_MACRO_CONNECTIONS": [  
  "gen_banks[0].u_bank vccd1 vssd1 vccd1 vssd1",  
  ...  
]
```

OpenLane reported "No match found" for each entry. Root cause: square brackets `[0]` are regex character classes, not literal brackets. The regex `gen_banks[0]` matches `gen_banksa`, `gen_banksb`, etc. — not the literal instance name.

Fix: Single wildcard entry `".*u_bank vccd1 vssd1 vccd1 vssd1"` matches all 8 bank instances correctly.

9.5 DRT_MIN_LAYER vs. RT_MIN_LAYER

After discovering that capacity adjustments rather than a hard layer ban were needed for SRAM macro pin access, the config used `DRT_MIN_LAYER=met3` to enforce signal routing to upper layers at detailed routing. This had no effect. Investigation of the OpenLane source (`set_routing_layers.tcl`) revealed: - `RT_MIN_LAYER` → sets minimum signal and clock routing layer (both GRT and DRT) - `DRT_MIN_LAYER` → only affects DRT clock nets; has no effect on signal routing

`DRT_MIN_LAYER` is the wrong variable and a common documentation trap.

9.6 PDN Stripe Collisions with Macro Row Placement

The initial top macro row at $y=750\text{ }\mu\text{m}$ caused 14 met4 routing DRC errors from a PDN horizontal power stripe that runs through the inter-row gap. Moving the top row to $y=850\text{ }\mu\text{m}$ cleared the stripe and resolved all met4 spacing violations.

Lesson: OpenLane's PDN stripe positions must be checked against macro placement before finalizing the floorplan; stripes cannot be routed through macro-to-macro gaps.

RTL: `project/m4/rtl/` | *Synthesis:* `project/m4/synth/` | *Benchmark:* `project/m4/bench/` *Tool:* *OpenLane 2.3.10*, *Yosys 0.46*, *OpenROAD*, *Sky130A sky130_fd_sc_hd* *Figures:* `report/figures/roofline_final.png` (*Figure 1 — roofline plot*)